


Lecture 3: Borrow checker

Barinov Denis

February 28, 2024

barinov.diu@gmail.com

Borrow Checker



What's the problem, Rust?

```
let mut v = vec![1, 2, 3];  
let x = &v[0];  
v.push(4);  
println!("{}", x);
```

What's the problem, Rust?

```
let mut v = vec![1, 2, 3];  
let x = &v[0];  
v.push(4);  
println!("{}", x);
```

error[E0502]: cannot borrow `v` as mutable because it is also borrowed as immutable

--> src/main.rs:8:5

```
|  
7 |     let x = &v[0];  
|           - immutable borrow occurs here  
8 |     v.push(4);  
|     ~~~~~ mutable borrow occurs here  
9 |     println!("{}", x);  
|           - immutable borrow later used here
```

What's the problem, Rust?

```
fn sum(v: Vec<i32>) -> i32 {  
    let mut result = 0;  
    for i in v {  
        result += i;  
    }  
    result  
}  
  
fn main() {  
    let mut v = vec![1, 2, 3];  
    println!("first sum: {}", sum(v));  
    v.push(4);  
    println!("second sum: {}", sum(v))  
}
```

What's the problem, Rust?

```
error[E0382]: borrow of moved value: `v`
  --> src/main.rs:12:5
    |
10 |     let mut v = vec![1, 2, 3];
    |         ----- move occurs because `v` has type `Vec<i32>`,
    |         which does not implement the `Copy` trait
11 |     println!("first sum: {}", sum(v));
    |                                   - value moved here
12 |     v.push(4);
    |     ^^^^^^^^^ value borrowed here after move
```

Ownership rules

- Each value in Rust has a variable that's called it's *owner*.
- There can be only one owner at a time.
- When the owner goes out of scope, the value will be dropped.

Ownership rules

```
fn main() {  
    let s = vec![1, 4, 8, 8];  
    let u = s;  
    println!("{:?}", u);  
    println!("{:?}", s); // This won't compile!  
}
```


Ownership rules

```
fn om_nom_nom(s: Vec<i32>) {  
    println!("I have consumed {s:?}");  
}  
  
fn main() {  
    let s = vec![1, 4, 8, 8];  
    om_nom_nom(s);  
    println!("{s:?}");  
}
```

Ownership rules

```
fn om_nom_nom(s: Vec<i32>) {  
    println!("I have consumed {s:?}");  
}
```

```
fn main() {  
    let s = vec![1, 4, 8, 8];  
    om_nom_nom(s);  
    println!("{s:?}");  
}
```

- Each "owner" has the responsibility to clean up after itself.
- When you move `s` into `om_nom_nom`, it becomes the owner of `s`, and it will free `s` when it's no longer needed in that scope. *Technically the `s` parameter in `om_nom_nom` become the owner.*
- That means you can no longer use it in `main`!
- In C++, we would create a copy!

Ownership rules

Given what we just saw, how can the following be the valid syntax?

```
fn om_nom_nom(n: u32) {  
    println!("{}", is a very nice number", n);  
}
```

```
fn main() {  
    let n: u32 = 42;  
    let m = n;  
    om_nom_nom(n);  
    om_nom_nom(m);  
    println!("{}", m + n);  
}
```

Ownership rules

- Say you have a group of lawyers that are reviewing and signing a contract over Google Docs (just pretend it's true :))
- What are some ground rules we'd need to set to avoid chaos?
- If someone modifies the contract before everyone else reviews/signs it, that's fine.
- But if someone modifies the contract while others are reviewing it, people might miss changes and think they're signing a contract that says something else.
- We should allow a single person to modify, or everyone to read, but not both.

Borrowing

- We can have multiple shared (immutable) references at once (with no mutable references) to a value.
- We can have only one mutable reference at once. (no shared references to it)
- This paradigm pops up a lot in systems programming, especially when you have "readers" and "writers". In fact, you've already studied it in the course of Theory and Practice of Concurrency.

Borrowing

- The lifetime of a value starts when it's created and ends the *last time it's used*
- Rust doesn't let you have a reference to a value that lasts longer than the value's lifetime
- Rust computes lifetimes at compile time using static analysis. (this is often an over-approximation!)
- Rust calls the special "drop" function on a value once its lifetime ends. (this is essentially a destructor)

Borrowing

```
fn main() {  
    let mut x = 5;  
    let y = &mut x;  
  
    println!("y = {y}");  
    x = 42; // ok  
    println!("x = {x}");  
}
```

Borrowing

```
fn main() {  
    let mut x = 5;  
    let y = &mut x;  
  
    x = 42; // not ok  
    println!("y = {y}");  
    println!("x = {x}");  
}
```


Borrowing

```
fn main() {  
    let x1 = 42;  
    let y1 = Box::new(84);  
    { // starts a new scope  
        let z = (x1, y1);  
        // z goes out of scope, and is dropped;  
        // it in turn drops the values from x1 and y1  
    }  
    // x1's value is Copy, so it was not moved into z  
    let x2 = x1;  
  
    // y1's value is not Copy, so it was moved into z  
    // let y2 = y1;  
}
```

Option¹ and Result²

Let's remember their definitions:

```
enum Option<T> {  
    Some(T),  
    None,  
}
```

```
enum Result<T, E> {  
    Ok(T),  
    Err(E),  
}
```

¹[Option documentation](#)

²[Result documentation](#)

Matching Option:

```
let result = Some("string");  
match result {  
    Some(s) => println!("String inside: {s}"),  
    None => println!("Ooops, no value"),  
}
```

Useful functions `.unwrap()` and `.expect()`:

```
fn unwrap(self) -> T;
```

```
fn expect(self, msg: &str) -> T;
```

Useful functions `.unwrap()` and `.expect()`:

```
let opt = Some(22022022);
assert!(opt.is_some());
assert!(!opt.is_none());
assert_eq!(opt.unwrap(), 22022022);
let x = opt.unwrap(); // Copy!

let newest_opt: Option<i32> = None;
// newest_opt.expect("I'll panic!");

let new_opt = Some(Vec::<i32>::new());
assert_eq!(new_opt.unwrap(), Vec::<i32>::new());
// error[E0382]: use of moved value: `new_opt`
// let x = new_opt.unwrap(); // Clone!
```

We have a magic function:

```
fn as_ref(&self) -> Option<&T>; // &self is &Option<T>
```

Let's solve a problem:

```
let new_opt = Some(Vec::<i32>::new());
assert_eq!(new_opt.unwrap(), Vec::<i32>::new());
// error[E0382]: use of moved value: `new_opt`
// let x = new_opt.unwrap(); // Clone!

let opt_ref = Some(Vec::<i32>::new());
assert_eq!(new_opt.as_ref().unwrap(), &Vec::<i32>::new());
let x = new_opt.unwrap(); // We used reference!
// There's also .as_mut() function
```

That means if type implements Copy, Option also implements Copy.

We can map `Option<T>` to `Option<U>`:

```
fn map<U, F>(self, f: F) -> Option<U>;
```

Example:

```
let maybe_some_string = Some(String::from("Hello, World!"));  
// `Option::map` takes self *by value*,  
// consuming `maybe_some_string`  
let maybe_some_len = maybe_some_string.map(|s| s.len());  
assert_eq!(maybe_some_len, Some(13));
```

There's **A LOT** of different `Option` functions, enabling us to write beautiful functional code:

```
fn map_or<U, F>(self, default: U, f: F) -> U;
fn map_or_else<U, D, F>(self, default: D, f: F) -> U;
fn unwrap_or(self, default: T) -> T;
fn unwrap_or_else<F>(self, f: F) -> T;
fn and<U>(self, optb: Option<U>) -> Option<U>;
fn and_then<U, F>(self, f: F) -> Option<U>;
fn or(self, optb: Option<T>) -> Option<T>;
fn or_else<F>(self, f: F) -> Option<T>;
fn xor(self, optb: Option<T>) -> Option<T>;
fn zip<U>(self, other: Option<U>) -> Option<(T, U)>;
```

It's recommended for you to study the documentation and try to avoid `match` where possible.

Option and ownership

There's two cool methods to control ownership of the value inside:

```
fn take(&mut self) -> Option<T>;  
fn replace(&mut self, value: T) -> Option<T>;  
fn insert(&mut self, value: T) -> &mut T;
```

The first one takes the value out of the `Option`, leaving a `None` in its place.

The second one replaces the value inside with the given one, returning `Option` of the old value.

The third one inserts a value into the `Option`, then returns a mutable reference to it.

Option API and ownership: take

```
struct Node<T> {  
    elem: T,  
    next: Option<Box<Node<T>>>,  
}  
  
pub struct List<T> {  
    head: Option<Box<Node<T>>>,  
}  
  
impl<T> List<T> {  
    pub fn pop(&mut self) -> Option<T> {  
        self.head.take().map(|node| {  
            self.head = node.next;  
            node.elem  
        })  
    }  
}
```

Rust guarantees to optimize the following types `T` such that `Option<T>` has the same size as `T`:

- `Box<T>`
- `&T`
- `&mut T`
- `fn`, `extern "C" fn`
- `#[repr(transparent)]` struct around one of the types in this list.
- `num::NonZero*`
- `ptr::NonNull<T>`

This is called the “null pointer optimization” or NPO.

Functions return `Result` whenever errors are expected and recoverable. In the `std` crate, `Result` is most prominently used for I/O.

Results must be used! A common problem with using return values to indicate errors is that it is easy to ignore the return value, thus failing to handle the error. `Result` is annotated with the `#[must_use]` attribute, which will cause the compiler to issue a warning when a `Result` value is ignored.³

³The Error Model

We can match it as a regular enum:

```
let version = Ok("1.1.14");  
match version {  
    Ok(v) => println!("working with version: {:?}", v),  
    Err(e) => println!("error: version empty"),  
}
```

We have pretty the same functionality as in Option:

```
fn is_ok(&self) -> bool;
fn is_err(&self) -> bool;
fn unwrap(self) -> T;
fn unwrap_err(self) -> E;
fn expect_err(self, msg: &str) -> E;
fn expect(self, msg: &str) -> T;
fn as_ref(&self) -> Result<&T, &E>;
fn as_mut(&mut self) -> Result<&mut T, &mut E>;
fn map<U, F>(self, op: F) -> Result<U, E>;
fn map_err<F, O>(self, op: O) -> Result<T, F>;
// And so on
```

It's recommended for you to study the documentation and try to avoid `match` where possible.

Operator ?

Consider the following structure:

```
struct Info {  
    name: String,  
    age: i32,  
}
```

Operator ?

```
fn write_info(info: &Info) -> io::Result<()> {  
    let mut file = match File::create("my_best_friends.txt") {  
        Err(e) => return Err(e),  
        Ok(f) => f,  
    };  
    if let Err(e) = file  
        .write_all(format!("name: {}\n", info.name)  
        .as_bytes()) {  
        return Err(e)  
    }  
    if let Err(e) = file  
        .write_all(format!("age: {}\n", info.age)  
        .as_bytes()) {  
        return Err(e)  
    }  
    Ok(())  
}
```



Operator ?

We can use the ? operator to make the code smaller!

```
fn write_info(info: &Info) -> io::Result<()> {  
    let mut file = File::create("my_best_friends.txt"?;  
    file.write_all(format!("name: {}\n", info.name).as_bytes())?;  
    file.write_all(format!("age: {}\n", info.age).as_bytes())?;  
    Ok(())  
}
```

Beautiful, isn't it?

We can use it for Option too!



Box **and** Rc

We are already familiar with Box type. Let's check one advanced function:

```
fn leak<'a>(b: Box<T, A>) -> &'a mut T;  
fn into_raw(b: Box<T, A>) -> *mut T;
```

Example:

```
let x = Box::new(41);  
let static_ref: &'static mut usize = Box::leak(x);  
*static_ref += 1;  
assert_eq!(*static_ref, 42);
```

But stop! Rust is the safe language, no memory unsafety, no undefined behavior, what's wrong!?

But stop! Rust is the safe language, no memory unsafety, no undefined behavior, what's wrong!?

*In reality, when you're creating global objects or interacting with other languages, you **have to** leak objects. Moreover, it's **safe** to leak memory, just not good!*

Rc is single-threaded reference-counting pointer. “Rc” stands for “Reference Counted”.

```
let rc = Rc::new(());  
let rc2 = rc.clone(); // Clones Rc, not what inside!  
let rc3 = Rc::clone(&rc); // The same
```

Rc is dropped when all instances of Rc are dropped.

Primary functions:

```
fn get_mut(this: &mut Rc<T>) -> Option<&mut T>;  
fn downgrade(this: &Rc<T>) -> Weak<T>;  
fn weak_count(this: &Rc<T>) -> usize;  
fn strong_count(this: &Rc<T>) -> usize;
```

References to the variable inside Rc are controlled at runtime:

```
let mut rc = Rc::new(42);  
println!("{}", *rc);  
  
*Rc::get_mut(&mut rc).unwrap() -= 41;  
println!("{}", *rc);  
  
let mut rc1 = rc.clone();  
println!("{}", *rc1);  
// thread 'main' panicked at 'called `Option::unwrap()`  
// on a `None` value'  
// *Rc::get_mut(&mut rc1).unwrap() -= 1;
```

`get_mut` guarantees that it will return mutable reference only if there's only one pointer. If there are more, you won't have a chance to modify Rc.

Rc is a **strong** pointer, while Weak is a **weak** pointer. Both of them have *ownership over allocation*, but only Rc have *ownership over the value inside*:

You can upgrade Weak to Rc:

```
fn upgrade(&self) -> Option<Rc<T>>;
```

```
let rc1 = Rc::new(String::from("string"));
let rc2 = rc1.clone();
let weak1 = Rc::downgrade(&rc1);
let weak2 = Rc::downgrade(&rc1);
drop(rc1); // The string is not deallocated
assert!(weak1.upgrade().is_some());
drop(weak1); // Nothing happens
drop(rc2); // The string is deallocated
assert_eq!(weak2.strong_count(), 0);
// If no strong pointers remain, this will return zero.
assert_eq!(weak2.weak_count(), 0);
assert!(weak2.upgrade().is_none());
drop(weak2); // The Rc is deallocated
```

There's also Arc - a thread-safe reference-counting pointer. Arc stands for "Atomically Reference Counted".

We'll need it to share data safely across threads in the future.